



## **Machine Learning Assignment**

### **PROJECT REPORT**

#### **TEAM 24**

#### **Detecting Payments Fraud Using Machine Learning**

<b>Name</b>	<b>SRN</b>
<b>Nagula Anish</b>	<b>PES2UG23CS358</b>
<b>Muskan Goenka</b>	<b>PES2UG23CS355</b>

## **Problem Statement**

Payments fraud is a rapidly growing threat, costing financial institutions and payment operators billions of dollars annually. Detecting fraudulent transactions in real time is difficult because fraudulent activity represents only a tiny fraction of total transactions, creating a severe class imbalance problem. Traditional detection methods often fail to identify rare but costly fraudulent activities while generating large numbers of false positives, impacting customer experience.

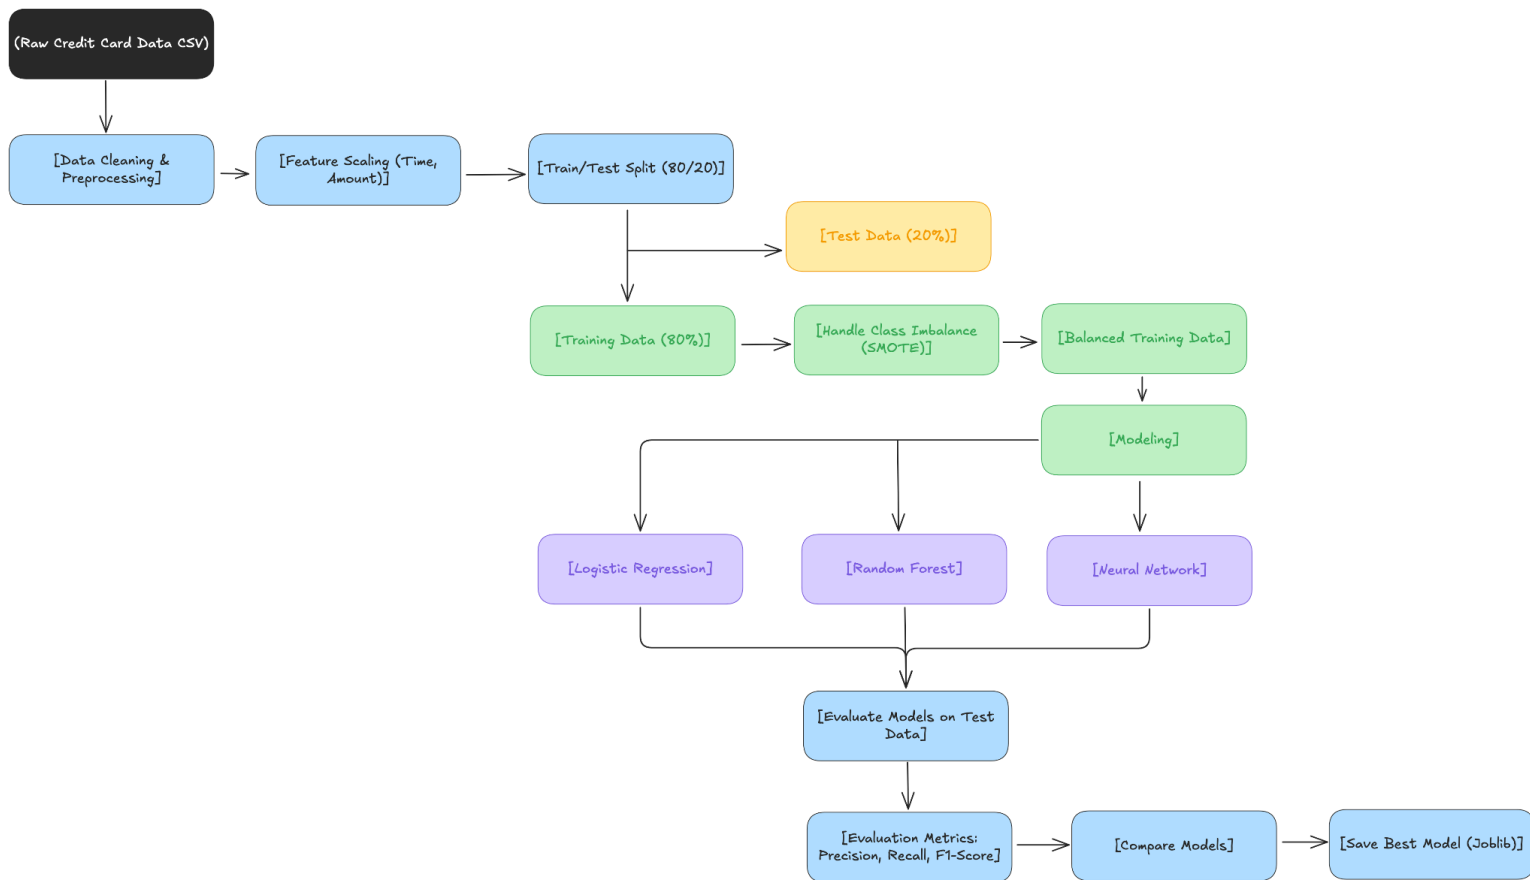
## **Objective / Aim**

- Develop machine learning models that can accurately distinguish between legitimate and fraudulent transactions in real-time.
- Address the challenge of severe class imbalance using advanced sampling techniques like SMOTE.
- Evaluate and compare different classification models using metrics suitable for imbalanced datasets such as Precision, Recall, and F1-Score.
- Deploy the best-performing model into an interactive web application for live demonstrations.

## **Dataset Details**

- Source: Kaggle Credit Card Fraud Detection Dataset
- Size: ~285,000 samples, 31 features
- Key Features:
  - **Time** ~ Seconds elapsed between each transaction and the first transaction.
  - **Amount** ~ The transaction amount.
  - **V1-V28** ~ Anonymized features resulting from a PCA transformation to protect user privacy.
- Target Variable: **Class** (0 -> Legitimate, 1 -> Fraudulent).

## Architecture Diagram



## Methodology

We followed a standard machine learning workflow for a classification problem, with the steps as outlined below:

- **Data Preparation:** Load the credit card fraud dataset using the Pandas library.
- **Feature Scaling:** Apply **StandardScaler** to the **Time** and **Amount** features to normalize their distributions and ensure they are on a comparable scale with the PCA-transformed **V** features.
- **Train/Test Split:** Split the data into an 80% training set and a 20% testing set. Stratification is used to ensure the proportion of fraudulent transactions is identical in both splits, which is crucial for achieving an unbiased evaluation.
- **Class Imbalance Handling:** Apply the *Synthetic Minority Over-sampling Technique* (SMOTE) only to the training set. This balances the class distribution by generating synthetic fraudulent samples, allowing the models to

learn their patterns without developing a bias towards the majority class which are the legitimate transactions in this case.

- **Model Training:** Train three distinct classification models on the balanced training data ~ Logistic Regression (as a baseline), Random Forest (an ensemble method), and a simple Neural Network.
- **Model Evaluation:** Evaluate the trained models on the original, unseen test set. Performance is measured using a classification report and confusion matrix for each model.
- **Persistence:** Save the final, best-performing model and the feature scaler to files using the `joblib` library for later use in a deployment environment.

## Results & Evaluation

The model performance is quantified using three key metrics ideal for imbalanced classification:

- **Precision:** Measures the accuracy of positive predictions. In this context, it answers: "Of all transactions flagged as fraud, what percentage were actually fraudulent?" A high precision minimizes false positives.
- **Recall (Sensitivity):** Measures the model's ability to find all positive samples. It answers: "Of all actual fraudulent transactions, what percentage did the model catch?" High recall is critical to minimize false negatives (missed fraud).
- **F1-Score:** The harmonic mean of Precision and Recall. It provides a single score that balances the two, which is useful for comparing overall model performance.

### Key Results

The evaluation of the three models on the test set yielded the following results for the **Fraud (Class 1)** category:

Metric	Logistic Regression	Random Forest	Neural Network	Interpretation
Precision	0.06 (6%)	<b>0.85 (85%)</b>	0.52 (52%)	The Random Forest model is highly trustworthy when it flags a transaction.
Recall	<b>0.92 (92%)</b>	0.84 (84%)	0.82 (82%)	Logistic Regression catches the

				highest percentage of fraud but at a great cost.
F1-Score	0.11	<b>0.84</b>	0.63	The Random Forest provides the best balance between precision and recall.

## Conclusion

The project successfully implemented and evaluated a machine learning pipeline to detect credit card fraud, confirming that model selection is critical when dealing with severely imbalanced data.

The results clearly show the trade-offs between different models. The **Logistic Regression** model, while achieving the highest recall of **0.92**, is impractical for real-world use due to its extremely low precision of **0.06**. This would lead to an unacceptably high number of legitimate customer transactions being incorrectly flagged as fraud.

The **Random Forest** classifier emerged as the clear winner. It achieved an excellent **F1-score of 0.84**, demonstrating a strong balance between a high precision (**0.85**) and a robust recall (**0.84**). This indicates that the model is both effective at catching fraud and reliable in its predictions.

### **Recommendation:**

The **Random Forest model is the recommended solution** for this fraud detection task. Its superior, well-balanced performance makes it the most suitable choice for deployment, as it effectively minimizes financial losses from fraud without causing significant disruption to legitimate customers.